Tree-in selection metrics across evidence-presentation conditions

Three models · six-node references · 100 planned tasks per model/condition · scores on a 0–1 scale

Model	Condition	Precision	Recall	Exact	Valid / N
Gemma-4-31B	K1: Source cards	1.000	0.972	0.870	100 / 100
	K3: Destination cards	1.000	0.983	0.920	100 / 100
	K4: Both endpoint cards	1.000	0.982	0.930	100 / 100
	K7: Direct query	1.000	0.982	0.940	100 / 100
Qwen3.6-Flash	K1: Source cards	0.988	0.603	0.150	100 / 100
	K3: Destination cards	1.000	0.603	0.160	100 / 100
	K4: Both endpoint cards	0.993	0.802	0.510	100 / 100
	K7: Direct query	0.978	0.660	0.160	100 / 100
Ministral-14B	K1: Source cards	0.976	0.853	0.540	100 / 100
	K3: Destination cards	0.986	0.787	0.380	99 / 100
	K4: Both endpoint cards	0.979	0.908	0.620	100 / 100
	K7: Direct query	0.875	0.837	0.620	91 / 100

K1/K3/K4 retain the procedural query and vary relation hosting; K4 repeats each relation on both endpoint cards.

K7 uses source-card relations and a direct query. Unique directed facts and targets match within each base task.

All means use planned denominators; invalid outcomes contribute zero. Gemma values are saved aggregate scores.

Source: archived Core, motif_size_metrics.csv. Descriptive values; no new confidence intervals or model calls.